

HW5: ML

Total points: 6.5

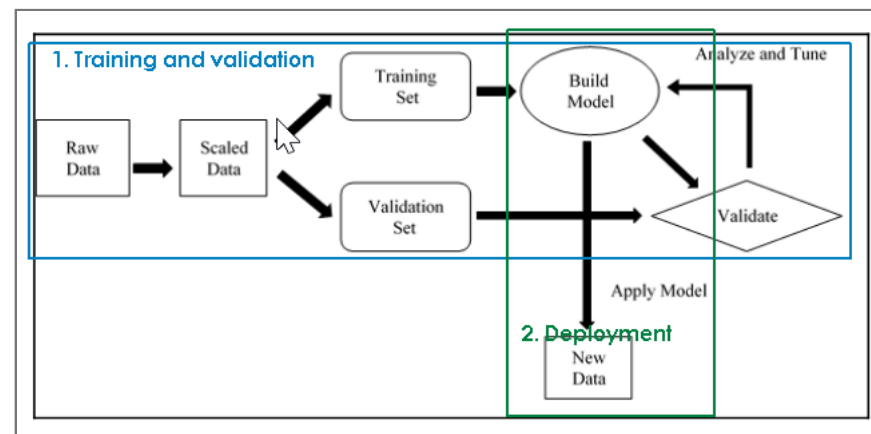
This last hw is on supervised **machine learning**! As you now know, it's **data-related** (lots, and lots, and lots of it), after all :)

Here is a summary of what you'll do:

- on Google's Colab (<https://colab.research.google.com/>), train a neural network on differentiating between a cat pic and dog pic, then use the trained network to classify a new (cat-like or dog-like) pic into a cat or dog
- after that, you'll train a tiny NN, and hand-verify that the training did work

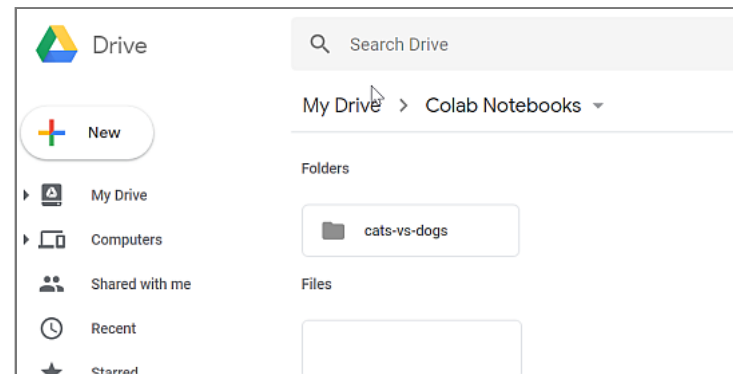
This is a 'soup-to-nuts' (start to finish) assignment that will get your feet wet (or plunge you in!), doing ML - a VERY valuable skill - training a self-driving car (<http://apollo.auto/>), for example, would involve much more complexity, but would be based on the same workflow.

You are going to carry out 'supervised learning', as shown in this annotated graphic [from a book on TensorFlow]:



Below are the steps. Have fun!

1. Use your GMail/GDrive account to log in, go to <https://drive.google.com/> (<https://drive.google.com/>), click on the '+ New' button at the top left of the page, look for the 'Colab' app [after + New, click on More >, then + Connect more apps] and connect it - this will make the app [which connects to the mighty Google Cloud on the other end!] be able to access (read, write) files and folders in your GDrive.
2. You'll notice that the above step created a folder called Colab Notebooks, inside your GDrive - this is good, because we can keep Colab-related things nicely organized inside that folder. Colab is a cloud environment (maintained by Google), for executing Jupyter 'notebooks'. A Jupyter notebook (.ipynb extension, 'Iron Python Notebook') is a JSON file that contains a mix of two types of "cells" - text cells that have Markdown-formatted text and images, and code cells that contain, well, code :) The code can be in **Julia**, **Python**, or **R** (or several other languages, including JavaScript, with appropriate language 'plugins' (kernels) installed); for this HW, we'll use Python notebooks. Here (nb/RealtimeR0.ipynb) is a COVID-19 notebook. Download it (make sure that ends up on your machine as a .ipynb extension; if your downloading turned it into a .txt, rename it to be .ipynb), then drag and drop it into your Colab GDrive folder. Colab will open the notebook - look through it, to see a mix of text, equations (can contain figures, videos... too), and, most significantly, code. The code in our notebook uses a portion of the data from <https://covidtracking.com/> (<https://covidtracking.com/>), to calculate 'R0' values. In Colab, do 'Runtime -> Run all' (pics/RealtimeR0.jpg) to run the code in all the cells; after the code is done running, scroll through, to see the results (several plots). **ALL the computations happened on the cloud, with data fetched by the code, from a URL** (https://bytes.usc.edu/~saty/data/2020_oct_nov_daily.csv (https://bytes.usc.edu/~saty/data/2020_oct_nov_daily.csv)) - pretty neat!
3. Within the Colab Notebooks subdir/folder, create a folder called cats-vs-dogs, for the hw:



Now we need DATA [images of cats and dogs] for training and validation, and scripts for training+validation and classifying.


4. Download this (data/data.zip) .zip data file (~85MB), unzip it. You'll see this structure:

```
data/
  live/
  train/
    cats/
    dogs/
  validation/
    cats/
    dogs/
```

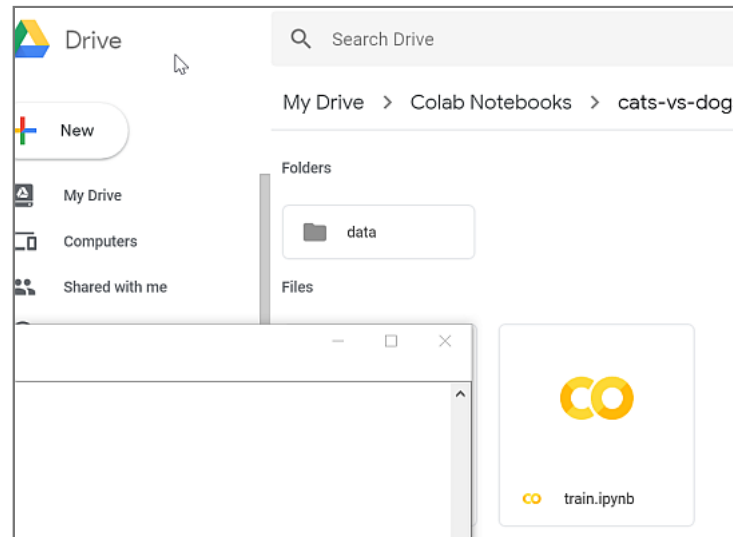
The train/ folder contains 1000 kittteh images under cats/, and 1000 doggo/pupper ones in dogs/. Have fun, looking at the adorable furballs :) Obviously **you** know which is which :) A neural network is going to start from scratch, and learn the difference, just based on these 2000 'training dataset' images. The validation/ folder contains 400 images each, of more cats and dogs - these are to feed the trained network, compare its classification answers to the actual answers so we can compute the accuracy of the training (in our code, we do this after each training epoch, to watch the accuracy build up, mostly monotonically). And, live/ is where you'd be placing new (to the NN) images of cats and dogs [that are not in the training or validation datasets], and use their filenames to **ask the network to classify them**: an output of 0 means 'cat', 1 means 'dog'. **Fun!**

Simply drag and drop the data/ folder on to your My Drive/Colab Notebooks/cats-vs-dogs/ area, and wait for about a half hour for the 2800 ($2 \times (1000 + 400)$) images to be uploaded. After that, you should be

seeing this [click inside the train/ and validation/ folders to see that the cats and dogs pics have been indeed uploaded]:



5. OK, **time to train a network!** Download this (nb/train.ipynb) Jupyter notebook [it is a '.ipynb' file, make sure you download it as such and not as .txt!]. Drag and drop the notebook into cats-vs-dogs/:



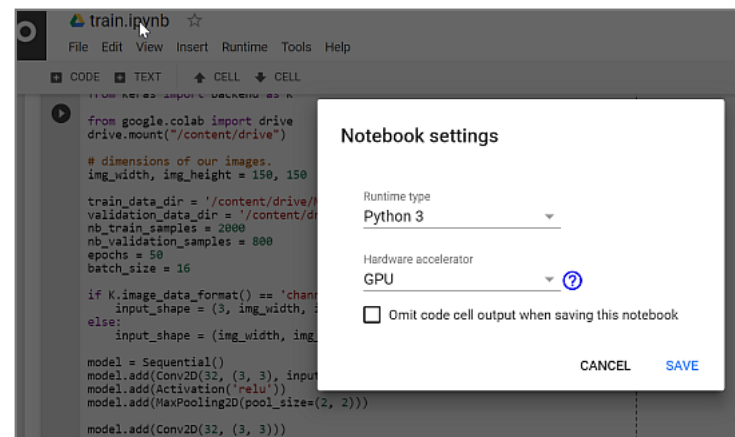
Double click on the notebook, that will open it so you can execute the code in the cell(s).

As you can see, it is a VERY short piece of code [not mine, except annotations and mods I made] where a network is set up [starting with 'model = Sequential()'], and the training is done using it [model.fit_generator()]. In the last line, the RESULTS [learned weights, biases, for each neuron in each layer] are stored on disk as a weights.h5

file [a .h5 file is binary, in the publicly documented .hd5 file format (https://en.wikipedia.org/wiki/Hierarchical_Data_Format) (hierarchical, JSON-like, perfect for storing network weights)].

The code uses the Keras NN library (<https://keras.io/>), which runs on graph (dataflow) execution backends such TensorFlow(TF), Theano, CNTK [here we are running it over TF via the Google cloud]. With Keras, it is possible to express NN architectures succinctly (<https://www.datacamp.com/community/blog/keras-cheat-sheet>) - the TF equivalent (or Theano's etc.) would be more verbose. As a future exercise, you can try coding the model in this hw, directly in TF or Theano or CNTK - you should get the same results.

Before you run the code to kick off the training, note that you will be using GPU acceleration on the cloud (**results in ~10x speedup**) - cool! You'd do this via 'Edit -> Notebook settings'. In this notebook, this is already set up (by me), but you can verify that it's set:

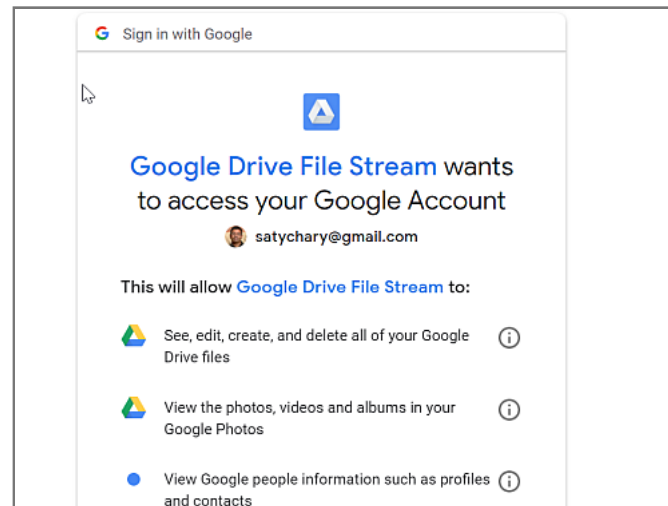



When you click on the circular 'play' button at the left of the cell, the training will start - here is a sped-up version of what you will get (your numerical values will be different):

0:00 / 0:11

The backprop loop runs 50 times ('epochs' (<https://keras.io/getting-started/faq/#what-does-sample-batch-epoch-mean>)) through all the training data. The acc: column shows the accuracy [how close the training is, to the expected validation/ results], which would be a little over 80% - NOT BAD, for having learned from just 1000 input images for each class!

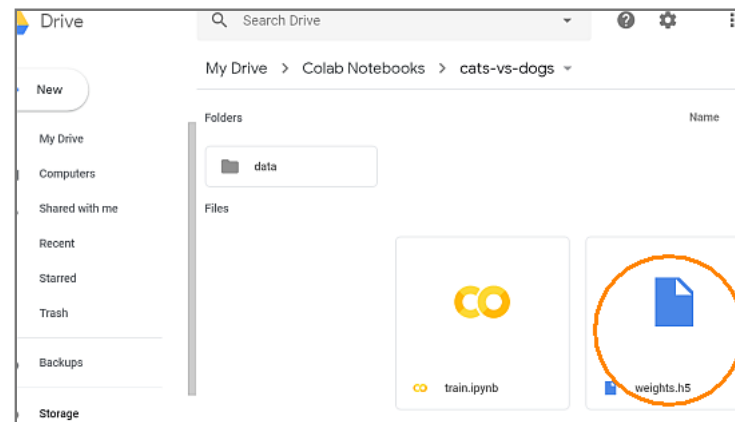
Click the play button to execute the code! The first time you run it (and anytime after logging out and logging back in), you'd need to authorize Colab to access GDrive - so a message will show up, under the code cell, asking you to click on a link whereby you can log in and provide authorization, and copy and paste the authorization code that appears. Once you do this, the rest of the code (where the training occurs) will start to run.



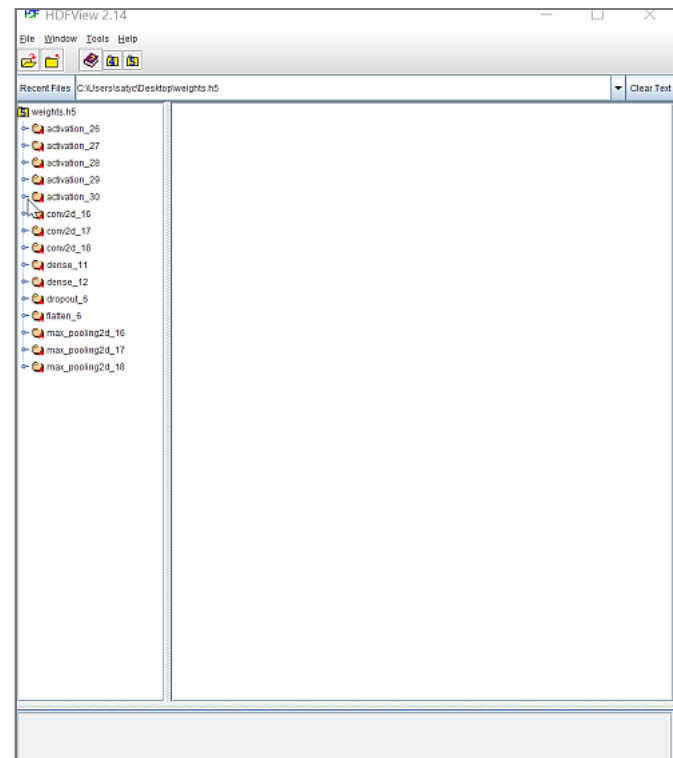


Scroll down to below the code cell, to watch the training happen. As you can see, it is going to take a short while.


After the 50th epoch, we're all done training (and validating too, which we did 50 times, once at the end of each epoch). **What's the tangible result, at the end of our training+validating process? It's a 'weights.h5' file!** If you look in your cats-vs-dogs/ folder, it should be there:




6. Soooo, what exactly [format and content-wise] is in the weights file? You can find out, by downloading HDFView-2.14.0, from <https://support.hdfgroup.org/products/java/release/download.html> (https://support.hdfgroup.org/products/java/release/download.html) [grab the binary, from the 'HDFView+Object 2.14' column on the left]. Install, and bring up the program. Download the .h5 file from GDrive to your local area (eg. desktop), then drag and drop it into HDView:



Right-click on weights.h5 at the top-left, and do 'Expand All':



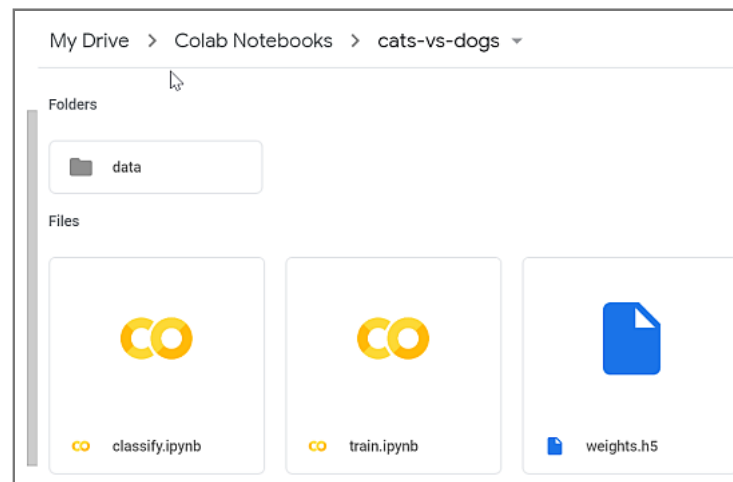
Neat! We can see the NN columns, and the biases and weights (kernels) for each. Double click on the bias and kernel items in the second (of the two) dense layers [dense_12, in my case - yours might be named something else], and stagger them so you can see both:



Computing those floating point numbers is **WHAT *EVERY* FORM OF NEURAL NETWORK TRAINING IS ALL ABOUT!** A self-driving car, for example, is also trained the same way, resulting in weights that can classify live traffic data (scary, in my opinion). Here, collectively (taking all layers into account), **it's those floating point numbers that REPRESENT the network's "learning" of telling apart cats and dogs!** The "learned" numbers (the .h5 weights file, actually) can be sent to anyone, who can instantiate a new network (with the same architecture as the one in the training step), and simply re/use the weights in weights.h5 to start classifying cats and dogs right away - no training necessary. The weight arrays represent "catness" and "dogness", in a sense :) We would call the network+weights, a 'pre-trained model'. In a self-driving car, the weights would be copied to the processing hardware (<https://www.wired.com/story/self-driving-cars-power-consumption-nvidia-chip/>) that resides in the car.

Q1 [0.5+0.5=1 point]. Submit your weights.h5 file. Also, create a submittable screengrab similar to the above [showing values for the second dense layer (eg. dense_12)]. For fun, click around, examine the arrays in the other layers as well. Again, it's all these values that are the end result of training, on account of iterating and minimizing classification errors through those epochs.

7. Now for the fun part - finding out how well our network has learned! Download this (nb/classify.ipynb) Jupyter notebook, and upload it to your GDrive's cats-vs-dogs/ Colab area [again, make sure the downloaded file is a .ipynb, NOT a .txt]:



When you open `classify.ipynb`, you can see that it contains Keras code to read the weights file and associate the weights with a new model (which needs to be 100% identical to the one we had set up, to train), then take a new image's filename as input, and **`predict(model.predict())` whether the image is that of a cat [output: 0], or a dog [output: 1]!** Why 0 for cat and 1 for dog? Because 'c' comes before 'd' alphabetically [or because (pics/purple.png)] :)

Supply (upload, into live/) a `what1.jpg` cat image, and `what2.jpg` dog image, then execute the cell. Hopefully you'd get a 0, and 1 (for `what1.jpg` and `what2.jpg`, respectively). The images can be any resolution (size) and aspect ratio (squarishness), but nearly-square pics

would work best. Try this with pics of your pets, your neighbors', images from a Google search, even your drawings/paintings... **Isn't this cool? Our little network can classify!**


Just FYI, note that the classification code in `classify.ipynb` could have simply been inside a new cell in `train.ipynb` instead. The advantage of multiple code cells inside a notebook, as opposed to multiple code blocks in a script, is that in a notebook, code cells can be independently executed one at a time (usually sequentially) - so if both of our programs were in the same notebook, we would run the training code first (just once), followed by classification (possibly multiple times); a script on the other hand, can't be re/executed in parts - that's because a Jupyter notebook is secretly a 'REPL' 'of sorts', it maintains 'state' [retains variables' values in a shared state used by all the cells] :)

Q2 [2 points]. Create a screenshot that shows the **[correct] classification** (you'll also be submitting your what{1,2}.jpg images with this).

What about misclassification? After all, we trained with "just" 1000 (not 1000000) images each, for about an 80% accurate prediction. What if we input 'difficult' images, of a cat that looks like it could be labeled a dog, and the other way around? :)

Q3 [2 points]. Get a 'Corgi' image [the world's smartest animals) dogs!], and a 'dog-like' cat image [hint, it's all about the ears!], upload to live/, **attempt to (mis)classify**, ie. create incorrect results (where the cat pic outputs a 1, and the dog's, 0), make a screenshot. Note that you need to edit the code to point `myPic` and `myPic2` to these image filenames. If you can't get a Corgi to be misclassified, you can use pics from other dog breeds with 'pointy' ears (<https://www.petside.com/dogs-with-pointy-ears/>), eg. Huskies :)

Q4 [1+0.5 = 1.5 points]. And now, on to something different. You are going to train the following neural network, and study its weights and outputs:



The leftmost layer is the input one, which will be 0 or 1. The three inputs go to the four neurons in the middle layer, and their outputs all go to the output neuron [which will output a 0 or 1].

Each neuron needs to 'learn' weights for its inputs, and a bias. In this case, we're making the bias to be 0 so we can just focus on the weights.

We'd like the neurons to learn this pattern [can YOU see what the pattern is? Hint: the X3 column is irrelevant, it's there just for the ride:]:

X1	X2	X3	Y
0	0	1	0
0	1	1	1
1	0	1	1
1	1	1	0

In other words, we want to find $Y = f(X1, X2, X3)$ that would produce the table above. $f()$ would contain 16 coefficients (weight values): 3 weights*4 = 12 weights for the middle layer, plus 4 weights for the last layer.

Create a new notebook, and enter this code:

```

import numpy as np

def sigmoid(x):
    return 1.0/(1+ np.exp(-x))

def sigmoid_derivative(x):
    return x * (1.0 - x)

class NeuralNetwork:
    def __init__(self, x, y):
        self.input      = x
        self.weights1   = np.random.rand(self.input.shape[1],4)
        self.weights2   = np.random.rand(4,1)
        self.y          = y
        self.output      = np.zeros(self.y.shape)

    def feedforward(self):
        self.layer1 = sigmoid(np.dot(self.input, self.weights1))
        self.output = sigmoid(np.dot(self.layer1, self.weights2))

    def backprop(self):
        # application of the chain rule to find derivative of the loss function with respect to weights2 and
        d_weights2 = np.dot(self.layer1.T, (2*(self.y - self.output) * sigmoid_derivative(self.output)))
        d_weights1 = np.dot(self.input.T,  (np.dot(2*(self.y - self.output) * sigmoid_derivative(self.output)
                                                    sigmoid_derivative(self.layer1)))

        # update the weights with the derivative (slope) of the loss function
        self.weights1 += d_weights1
        self.weights2 += d_weights2

if __name__ == "__main__":
    x = np.array([[0,0,1],
                  [0,1,1],
                  [1,0,1],
                  [1,1,1]])
    y = np.array([[0],[1],[1],[0]])
    nn = NeuralNetwork(x,y)

    for i in range(10000):
        nn.feedforward()
        nn.backprop()

    print(nn.weights1) # 'blue' layer  [hidden]
    print(nn.weights2) # 'green' layer [output]
    nn.feedforward()   # this was missing earlier, thanks to Man Sun for noticing!
    print(nn.output)   # predictions

```

You can study the code (optionally) to see what it does - it does NOT use any ML library (eg TF, Keras, PyTorch...) at all, as you can see - it is all coded from first principles, using just numpy [which we use, for vector ops; it's possible to code these ops too from scratch, and avoid using numpy as well, making it even more of a pure, standalone implementation].

Run the code, and look at the weights, and the output. Eg. here is a sample run [note: your weights will have different values compared to mine]:

```


[ [ 7.19328546  2.81204434  1.40998559  4.83853315 ]
  [ 6.35378132  4.51126787  3.38229611 -3.55865088 ]
  [-3.03899765 -5.5302491  -3.51765491  2.34739304 ] ]
[ [11.72259331]
  [-8.30218083]
  [-5.14836988]
  [-5.71845324] ]
-----
[0.00765428]
[0.9954903 ]
[0.99146861]
[0.00676557]]

```

As you can see from nn.output [last 4 printed values above], the values that are supposed to be 0,1,1,0, are pretty close to them. So the matrices of top 16 floating numbers represent "intelligence", they are the result of the NN having "learned" the **XOR truth table** [THAT is the answer to the 'what is this pattern' question above] :)

Q4a. On paper, write out a **SINGLE equation $f(X_1, X_2, X_3)$** that shows the output **Y** in terms of the three variables **X1, X2 and X3**, and the **neurons' weights!** We can number the three neural layers (input, hidden, output) 0,1,2, and in each column, number the neurons from top to bottom, 0,1... So every neuron's weight w will have three subscripts: first one for weight# (0,1...), second one for column #, third one for neuron # within the column.

Here are notes to provide more clarity on the above. The weights numbering indicated above, is simply for you to keep track of what weight (from your printed output) goes where (these numberings will not be in your equation that you'd submit). As an example, here is the numbering for the three weights, #0,#1,#2, of just one neuron (which is highlighted in the left pic: column#1, neuron#2):



The equation that you are asked to create, will not contain coeffs such as $w_{2,1,2}$ etc - instead, it would contain the **actual values** for them, eg, -0.8234.

Q4b. Use the equation you just created, to **hand-calculate the classification** (on paper!). In other words, when 0,0,1 is the input, use the weights to process this input, and show that the output would be 0.00765428, to use my output above as an example [likewise the other three rows should produce the three output values shown above]. Remember to pass the summed weights through the sigmoid (non-linear) function [look at the code, for the sigmoid formula that involves exponentiation]! **Again - do this for all the four input triples** [generate the output for 001, 011, 101, 111]. The point of this is to make you see this: once the weights have been learned, it's a straightforward process to handle incoming data to generate outputs [it is not complex/mysterious!], **it simply involves evaluating the RHS of a deterministic equation**. In this exercise, we're simply using the training data and the learned weights, to hand calculate the outputs; in much more complex situations, the network will be fed new (so far unseen) data, which it will process using learned weights (as if through a formula!), so the overall workflow/idea is the same.

Here's a checklist of what to submit **[as a single .zip file]**:

- weights.h5, and a screenshot from HDFView
 - your 'good' cat and dog pics, and screenshot that shows proper classification
 - your 'trick' cat and dog pics, and screenshot that shows misclassification
 - pics of the NN equation, and the four calculations (one for each row of data)
-

No-points bonus. Add a third class to the first network above, eg. racoon (https://www.google.com/search?q=racoon&rlz=1C1CHBF_enUS723US723&sxsrf=ACYBGNQEveXxDEn19VOiT6UfOQg94Wpe1Q:15748850523

to the training. You'd need to create (acquire) (pics/img_scraping.jpg) your own training data (500 images might be sufficient, or, even fewer,

eg. 300), and **modify the code** just a tiny bit in order to train, and test, on this third class, too. The third class can even be a non-animal, eg. trash (cardboard or plastic or metal or glass), ball, furniture, book... anything at all! Here ([nb/train3Classes.ipynb](#)) is the training script, and here ([nb/classify3Classes.ipynb](#)) is the classifying script.

All done - hope you had fun, and learned a lot!

Note - you can continue using Colab to run all sorts of notebooks (<https://github.com/jupyter/jupyter/wiki/A-gallery-of-interesting-Jupyter-Notebooks>) [on Google's cloud GPUs!], including ones with TensorFlow, Keras, PyTorch... etc. ML code.
